

Understanding Java Garbage Collection

and what you can do about it

Gil Tene, CTO & co-Founder, Azul Systems



This Talk's Purpose / Goals

- ☉ This talk is focused on GC education
- ☉ This is not a “how to use flags to tune a collector” talk
- ☉ This is a talk about how the “GC machine” works
- ☉ Purpose: Once you understand how it works, you can use your own brain...
- ☉ You'll learn just enough to be dangerous...
- ☉ The “Azul makes the world's greatest GC” stuff will only come at the end, I promise...

About me: Gil Tene

- co-founder, CTO
@Azul Systems
- Have been working on a “think different” GC approaches since 2002
- Created Pauseless & C4 core GC algorithms
(Tene, Wolf)
- A Long history building Virtual & Physical Machines, Operating Systems, Enterprise apps, etc...





* working on real-world trash compaction issues, circa 2004

About Azul

- ☛ We make scalable Virtual Machines
- ☛ Have built “whatever it takes to get job done” since 2002
- ☛ 3 generations of custom SMP Multi-core HW (Vega)
- ☛ **Now Pure software for commodity x86 (Zing)**
- ☛ “Industry firsts” in Garbage collection, elastic memory, Java virtualization, memory scale

Zing



High level agenda

- ④ GC fundamentals and key mechanisms
- ④ Some GC terminology & metrics
- ④ Classifying currently available collectors
- ④ The “Application Memory Wall” problem
- ④ The C4 collector: What an actual solution looks like...

Memory use

How many of you use heap sizes of:



more than 1/2 GB?



more than 1 GB?



more than 2 GB?



more than 4 GB?



more than 10 GB?



more than 20 GB?



more than 50 GB?

Why should you care?

The story of the good little architect

- ④ A good architect must, first and foremost, be able to impose their architectural choices on the project...
- ④ Early in Azul's concurrent collector days, we encountered an application exhibiting 18 second pauses
 - ④ Upon investigation, we found the collector was performing 10s of millions of object finalizations per GC cycle
 - *We have since made reference processing fully concurrent...
- ④ Every single class written in the project has a finalizer
 - ④ The only work the finalizers did was nulling every reference field
- ④ The right discipline for a C++ ref-counting environment
 - ④ The wrong discipline for a precise garbage collected environment

Trying to solve GC problems in application architecture is live throwing knives

- ☹ You probably shouldn't do it blindfolded
- ☹ It takes practice and understanding to get it right
- ☹ You can get very good at it, but do you really want to?
 - ☹ Will all the code you leverage be as good as yours?
- ☹ Examples:
 - ☹ Object pooling
 - ☹ Off heap storage
 - ☹ Distributed heaps
 - ☹ ...
 - ☹ (In most cases, you end up building your own garbage collector)

Most of what People seem to “know” about Garbage Collection is wrong

- ④ In many cases, it's much better than you may think
 - ④ GC is extremely efficient. Much more so than malloc()
 - ④ Dead objects cost nothing to collect
 - ④ GC will find all the dead objects (including cyclic graphs)
 - ④ ...
- ④ In many cases, it's much worse than you may think
 - ④ Yes, it really does stop for ~1 sec per live GB.
 - ④ No, GC does not mean you can't have memory leaks
 - ④ No, those pauses you eliminated from your 20 minute test are not gone
 - ④ ...

Some GC Terminology

A Basic Terminology example:

What is a concurrent collector?

- ☉ A Concurrent Collector performs garbage collection work concurrently with the application's own execution
- ☉ A Parallel Collector uses multiple CPUs to perform garbage collection

Classifying a collector's operation

- ④ A Concurrent Collector performs garbage collection work concurrently with the application's own execution
- ④ A Parallel Collector uses multiple CPUs to perform garbage collection
- ④ A Stop-the-World collector performs garbage collection while the application is completely stopped
- ④ An Incremental collector performs a garbage collection operation or phase as a series of smaller discrete operations with (potentially long) gaps in between
- ④ Mostly means sometimes it isn't (usually means a different fall back mechanism exists)

Precise vs. Conservative Collection

- ④ A Collector is Conservative if it is unaware of all object references at collection time, or is unsure about whether a field is a reference or not
- ④ A Collector is Precise if it can fully identify and process all object references at the time of collection
 - ④ A collector **MUST** be precise in order to move objects
 - ④ The COMPILERS need to produce a lot of information (oopmaps)
- ④ All commercial server JVMs use precise collectors
 - ④ All commercial server JVMs use some form of a moving collector

Safepoints

- ③ A GC Safepoint is a point or range in a thread's execution where the collector can identify all the references in that thread's execution stack
 - ③ "Safepoint" and "GC Safepoint" are often used interchangeably
 - ③ But there are other types of safepoints, including ones that require more information than a GC safepoint does (e.g. deoptimization)
- ③ "Bringing a thread to a safepoint" is the act of getting a thread to reach a safepoint and not execute past it
 - ③ Close to, but not exactly the same as "stop at a safepoint"
 - ③ e.g. JNI: you can keep running in, but not past the safepoint
 - ③ Safepoint opportunities are (or should be) frequent
- ③ In a Global Safepoint all threads are at a Safepoint

What's common to all precise GC mechanisms?

- ④ Identify the live objects in the memory heap
- ④ Reclaim resources held by dead objects
- ④ Periodically relocate live objects
- ④ Examples:
 - ④ Mark/Sweep/Compact (common for Old Generations)
 - ④ Copying collector (common for Young Generations)

Mark (aka "Trace")

- ④ Start from "roots" (thread stacks, statics, etc.)
- ④ "Paint" anything you can reach as "live"
- ④ At the end of a mark pass:
 - ④ all reachable objects will be marked "live"
 - ④ all non-reachable objects will be marked "dead" (aka "non-live").
- ④ Note: work is generally linear to "live set"

Sweep

- ④ Scan through the heap, identify “dead” objects and track them somehow
 - ④ (usually in some form of free list)
- ④ Note: work is generally linear to heap size

Compact

- ④ Over time, heap will get “swiss cheesed”: contiguous dead space between objects may not be large enough to fit new objects (aka “fragmentation”)
- ④ Compaction moves live objects together to reclaim contiguous empty space (aka “relocate”)
- ④ Compaction has to correct all object references to point to new object locations (aka “remap”)
- ④ Remap scan must cover all references that could possibly point to relocated objects
- ④ Note: work is generally linear to “live set”

Copy

- Copying collector moves all live objects from a “from” space to a “to” space & reclaims “from” space
- At start of copy, all objects are in “from” space and all references point to “from” space.
- Start from “root” references, copy any reachable object to “to” space, correcting references as we go
- At End of copy, all objects are in “to” space, and all references point to “to” space
- Note: work generally linear to “live set”

Mark/Sweep/Compact, Copy, Mark/Compact

- Copy requires 2x the max. live set to be reliable
- Mark/Compact [typically] requires 2x the max. live set in order to fully recover garbage in each cycle
- Mark/Sweep/Compact only requires 1x (plus some)
- Copy and Mark/Compact are linear only to live set
- Mark/Sweep/Compact linear (in sweep) to heap size
- Mark/Sweep/(Compact) may be able to avoid some moving work
- Copying is [typically] “monolithic”

Generational Collection

- ④ Generational Hypothesis: most objects die young
- ④ Focus collection efforts on young generation:
 - ④ Use a moving collector: work is linear to the live set
 - ④ The live set in the young generation is a small % of the space
 - ④ Promote objects that live long enough to older generations
- ④ Only collect older generations as they fill up
 - ④ “Generational filter” reduces rate of allocation into older generations
- ④ Tends to be (order of magnitude) more efficient
 - ④ Great way to keep up with high allocation rate
 - ④ Practical necessity for keeping up with processor throughput

Generational Collection

- ④ Requires a “Remembered set”: a way to track all references into the young generation from the outside
- ④ Remembered set is also part of “roots” for young generation collection
- ④ No need for 2x the live set: Can “spill over” to old gen
- ④ Usually want to keep surviving objects in young generation for a while before promoting them to the old generation
 - ④ Immediate promotion can dramatically reduce gen. filter efficiency
 - ④ Waiting too long to promote can dramatically increase copying work

How does the remembered set work?

- ④ Generational collectors require a “Remembered set”: a way to track all references into the young generation from the outside
- ④ Each store of a NewGen reference into an OldGen object needs to be intercepted and tracked
- ④ Common technique: “Card Marking”
 - ④ A bit (or byte) indicating a word (or region) in OldGen is “suspect”
- ④ Write barrier used to track references
 - ④ Common technique (e.g. HotSpot): blind stores on reference write
 - ④ Variants: precise vs. imprecise card marking, conditional vs. non-conditional

The typical combos in commercial server JVMs

- ④ Young generation usually uses a copying collector
- ④ Young generation is usually monolithic, stop-the-world
- ④ Old generation usually uses Mark/Sweep/Compact
- ④ Old generation may be STW, or Concurrent, or mostly-Concurrent, or Incremental-STW, or mostly-Incremental-STW

Useful terms for discussing garbage collection

☉ Mutator

- ☉ Your program...

☉ Parallel

- ☉ Can use multiple CPUs

☉ Concurrent

- ☉ Runs concurrently with program

☉ Pause

- ☉ A time duration in which the mutator is not running any code

☉ Stop-The-World (STW)

- ☉ Something that is done in a pause

☉ Monolithic Stop-The-World

- ☉ Something that must be done in its entirety in a single pause

☉ Generational

- ☉ Collects young objects and long lived objects separately.

☉ Promotion

- ☉ Allocation into old generation

☉ Marking

- ☉ Finding all live objects

☉ Sweeping

- ☉ Locating the dead objects

☉ Compaction

- ☉ Defragments heap
- ☉ Moves objects in memory
- ☉ Remaps all affected references
- ☉ Frees contiguous memory regions

Useful metrics for discussing garbage collection

● Heap population (aka Live set)

- How much of your heap is alive

● Allocation rate

- How fast you allocate

● Mutation rate

- How fast your program updates references in memory

● Heap Shape

- The shape of the live object graph
- * Hard to quantify as a metric...

● Object Lifetime

- How long objects live

● Cycle time

- How long it takes the collector to free up memory

● Marking time

- How long it takes the collector to find all live objects

● Sweep time

- How long it takes to locate dead objects
- * Relevant for Mark-Sweep

● Compaction time

- How long it takes to free up memory by relocating objects
- * Relevant for Mark-Compact

Empty memory and CPU/throughput

Two Intuitive limits

- ☉ If we had infinite empty memory, we would never have to collect, and GC would take 0% of the CPU time
- ☉ If we had exactly 1 byte of empty memory at all times, the collector would have to work “very hard”, and GC would take 100% of the CPU time
- ☉ GC CPU % will follow a rough $1/x$ curve between these two limit points, dropping as the amount of memory increases.

Empty memory needs

(empty memory == CPU power)

- ④ The amount of empty memory in the heap is the dominant factor controlling the amount of GC work
- ④ For both Copy and Mark/Compact collectors, the amount of work per cycle is linear to live set
- ④ The amount of memory recovered per cycle is equal to the amount of unused memory (heap size) - (live set)
- ④ The collector has to perform a GC cycle when the empty memory runs out
- ④ A Copy or Mark/Compact collector's efficiency doubles with every doubling of the empty memory

What empty memory controls

- ④ Empty memory controls efficiency (amount of collector work needed per amount of application work performed)
- ④ Empty memory controls the frequency of pauses (if the collector performs any Stop-the-world operations)
- ④ Empty memory DOES NOT control pause times (only their frequency)
- ④ In Mark/Sweep/Compact collectors that pause for sweeping, more empty memory means less frequent but LARGER pauses

Some non-monolithic-STW stuff

Concurrent Marking

- ④ Mark all reachable objects as “live”, but object graph is “mutating” under us.
- ④ Classic concurrent marking race: mutator may move reference that has not yet been seen by the marker into an object that has already been visited
 - ④ If not intercepted or prevented in some way, will corrupt the heap
- ④ Example technique: track mutations, multi-pass marking
 - ④ Track reference mutations during mark (e.g. in card table)
 - ④ Re-visit all mutated references (and track new mutations)
 - ④ When set is “small enough”, do a STW catch up (mostly concurrent)
- ④ Note: work grows with mutation rate, may fail to finish

Incremental Compaction

- ④ Track cross-region remembered sets (which region points to which)
- ④ To compact a single region, only need to scan regions that point into it to remap all potential references
- ④ identify regions sets that fit in limited time
 - ④ Each such set of regions is a Stop-the-World increment
 - ④ Safe to run application between (but not within) increments
- ④ Note: work can grow with the square of the heap size
 - ④ The number of regions pointing into a single region is generally linear to the heap size (the number of regions in the heap)

Delaying the inevitable

- Compaction is inevitable in practice
 - And compacting anything requires scanning/fixing all references to it
- Delay tactics focus on getting “easy empty space” first
 - This is the focus for the vast majority of GC tuning
- Most objects die young [Generational]
 - So collect young objects only, as much as possible
 - But eventually, some old dead objects must be reclaimed
- Most old dead space can be reclaimed without moving it
 - [e.g. CMS] track dead space in lists, and reuse it in place
 - But eventually, space gets fragmented, and needs to be moved
- Much of the heap is not “popular” [e.g. G1, “Balanced”]
 - A non popular region will only be pointed to from a small % of the heap
 - So compact non-popular regions in short stop-the-world pauses
 - But eventually, popular objects and regions need to be compacted

Classifying common collectors

The typical combos in commercial server JVMs

- ④ Young generation usually uses a copying collector
 - ④ Young generation is usually monolithic, stop-the-world
- ④ Old generation usually uses a Mark/Sweep/Compact collector
 - ④ Old generation may be STW, or Concurrent, or mostly-Concurrent, or Incremental-STW, or mostly-Incremental-STW

HotSpot™ ParallelGC

Collector mechanism classification

- ④ Monolithic Stop-the-world copying NewGen
- ④ Monolithic Stop-the-world Mark/Sweep/Compact OldGen

HotSpot™ ConcMarkSweepGC (aka CMS)

Collector mechanism classification

- ④ Monolithic Stop-the-world copying NewGen (ParNew)
- ④ Mostly Concurrent, non-compacting OldGen (CMS)
 - ④ Mostly Concurrent marking
 - ④ Mark concurrently while mutator is running
 - ④ Track mutations in card marks
 - ④ Revisit mutated cards (repeat as needed)
 - ④ Stop-the-world to catch up on mutations, ref processing, etc.
 - ④ Concurrent Sweeping
 - ④ Does not Compact (maintains free list, does not move objects)
- ④ Fallback to Full Collection (Monolithic Stop the world).
 - ④ Used for Compaction, etc.

HotSpot™ G1GC (aka “Garbage First”)

Collector mechanism classification

- ④ Monolithic Stop-the-world copying NewGen
- ④ Mostly Concurrent, OldGen marker
 - ④ Mostly Concurrent marking
 - ④ Stop-the-world to catch up on mutations, ref processing, etc.
 - ④ Tracks inter-region relationships in remembered sets
- ④ Stop-the-world mostly incremental compacting old gen
 - ④ Objective: “Avoid, as much as possible, having a Full GC...”
 - ④ Compact sets of regions that can be scanned in limited time
 - ④ Delay compaction of popular objects, popular regions
- ④ Fallback to Full Collection (Monolithic Stop the world).
 - ④ Used for compacting popular objects, popular regions, etc.

The “Application Memory Wall”

Memory use

How many of you use heap sizes of:



more than 1/2 GB?



more than 1 GB?



more than 2 GB?



more than 4 GB?



more than 10 GB?



more than 20 GB?



more than 50 GB?

Reality check: servers in 2011

- ☉ Retail prices, major web server store (US \$, Oct. 2011)

24 vCore, 96GB server \approx \$5K

32 vCore, 256GB server \approx \$16K

64 vCore, 512GB server \approx \$30K

80 vCore, 1TB server \approx \$63K

- ☉ Cheap (\approx \$1.5/GB/Month), and roughly linear to \sim 1TB

- ☉ 10s to 100s of GB/sec of memory bandwidth

The Application Memory Wall

A simple observation:

- ④ Application instances appear to be unable to make effective use of modern server memory capacities
- ④ The size of application instances as a % of a server's capacity is rapidly dropping

How much memory do applications need?

- “640KB ought to be enough for anybody”

“I've said some stupid things and some wrong things, but not that. No one involved in computers would ever say that a certain amount of memory is enough for all time ...” - Bill Gates, 1996

WRONG!

- So what's the right number?

6,400K?

64,000K?

640,000K?

6,400,000K?

64,000,000K?

- There is no right number
- Target moves at 50x-100x per decade

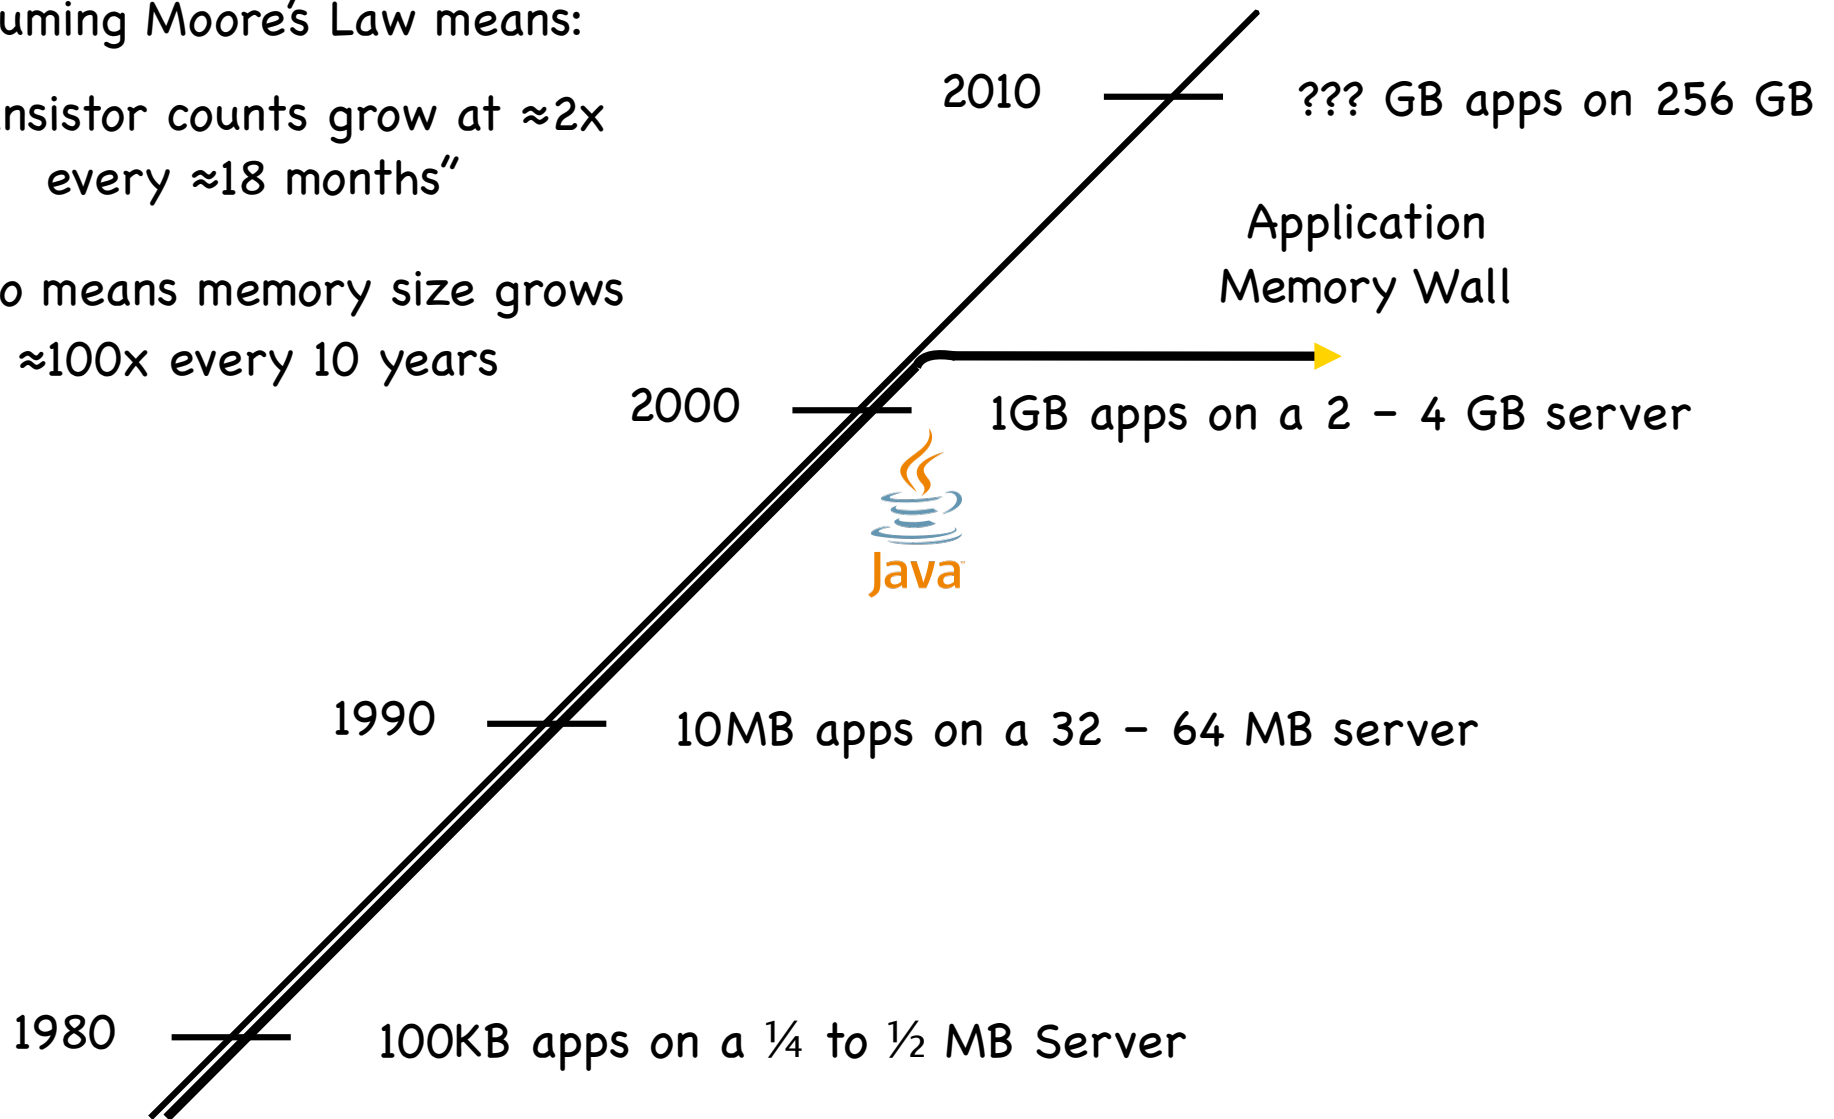


"Tiny" application history

Assuming Moore's Law means:

"transistor counts grow at $\approx 2\times$
every ≈ 18 months"

It also means memory size grows
 $\approx 100\times$ every 10 years



* "Tiny": would be "silly" to distribute

What is causing the Application Memory Wall?

- Garbage Collection is a clear and dominant cause
- There seem to be practical heap size limits for applications with responsiveness requirements
- [Virtually] All current commercial JVMs will exhibit a multi-second pause on a normally utilized 2-4GB heap.
 - It's a question of "When" and "How often", not "If".
 - GC tuning only moves the "when" and the "how often" around
- Root cause: The link between scale and responsiveness
 - Without removing link, can't use >~\$500 worth of 2011 H/W.
 - Problem/gap growing at Moore's law rate

What quality of GC is responsible for the Application Memory Wall?

- ④ It is NOT about overhead or efficiency:
 - ④ CPU utilization, bottlenecks, memory consumption and utilization
- ④ It is NOT about speed
 - ④ Average speeds, 90%, 99% speeds, are all perfectly fine
- ④ It is NOT about minor GC events (right now)
 - ④ GC events in the 10s of msec are usually tolerable for most apps
- ④ It is NOT about the frequency of very large pauses
- ④ It is all about the worst observable pause behavior
 - ④ People avoid building/deploying visibly broken systems

GC Problems

Framing the discussion:

Garbage Collection at modern server scales

- ④ Modern Servers have 100s of GB of memory
- ④ Each modern x86 core (when actually used) produces garbage at a rate of $\frac{1}{4} - \frac{1}{2}$ GB/sec +
- ④ That's many GB/sec of allocation in a server
- ④ Monolithic stop-the-world operations are the cause of the current Application Memory Wall
 - ④ Even if they are done "only a few times a day"

The things that seem “hard” to do in GC

④ Robust concurrent marking

- ④ References keep changing
- ④ Multi-pass marking is sensitive to mutation rate
- ④ Weak, Soft, Final references “hard” to deal with concurrently

④ [Concurrent] Compaction...

- ④ It's not the moving of the objects...
- ④ It's the fixing of all those references that point to them
- ④ How do you deal with a mutator looking at a stale reference?
- ④ If you can't, then remapping is a [monolithic] STW operation

④ Young Generation collection at scale

- ④ Young Generation collection is generally monolithic, Stop-The-World
- ④ Young generation pauses are only small because heaps are tiny
- ④ A 100GB heap will regularly have several GB of live young stuff...

How can we break through the
Application Memory Wall?

We need to solve the right problems

- ④ Focus on the causes of the Application Memory Wall
 - ④ Scale is artificially limited by responsiveness
- ④ Responsiveness must be unlinked from scale:
 - ④ Heap size, Live Set size, Allocation rate, Mutation rate
 - ④ Responsiveness must be continually sustainable
 - ④ Can't ignore "rare" events
- ④ Eliminate all Stop-The-World Fallbacks
 - ④ At modern server scales, any STW fall back is a failure

The problems that need solving

(areas where the state of the art needs improvement)

④ Robust Concurrent Marking

- ④ In the presence of high mutation and allocation rates
- ④ Cover modern runtime semantics (e.g. weak refs)

④ Compaction that is not monolithic-stop-the-world

- ④ Stay responsive while compacting $\frac{1}{4}$ TB heaps
- ④ Must be robust: not just a tactic to delay STW compaction
- ④ [current “incremental STW” attempts fall short on robustness]

④ Young-Gen that is not monolithic-stop-the-world


- ④ Stay responsive while promoting multi-GB data spikes
- ④ Concurrent or “incremental STW” may be both be ok
- ④ Surprisingly little work done in this specific area

Azul's "C4" Collector

Continuously Concurrent Compacting Collector

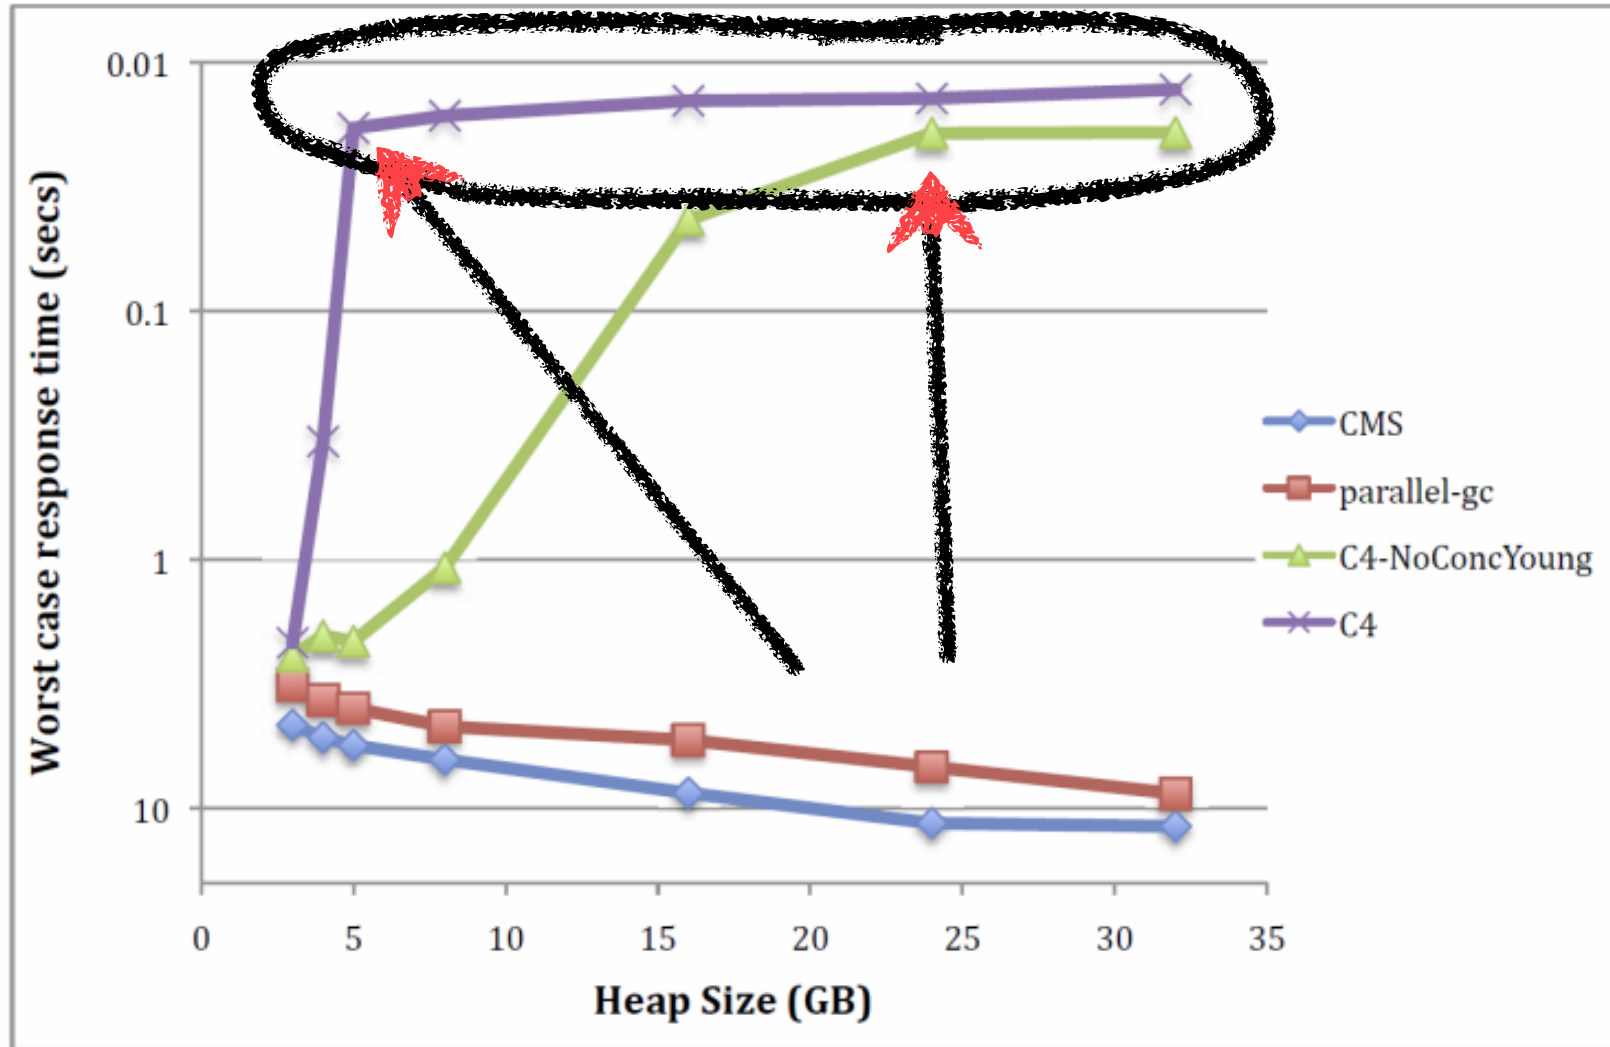
- ④ Concurrent, compacting new generation
- ④ Concurrent, compacting old generation
- ④ Concurrent guaranteed-single-pass marker
 - ④ Oblivious to mutation rate
 - ④ Concurrent ref (weak, soft, final) processing
- ④ Concurrent Compactor
 - ④ Objects moved without stopping mutator
 - ④ References remapped without stopping mutator
 - ④ Can relocate entire generation (New, Old) in every GC cycle
- ④ No stop-the-world fallback
 - ④ Always compacts, and always does so concurrently

Sample throughput



- SpecJBB + Slow churning 2GB LRU Cache
- Live set is ~2.5GB across all measurements
- Allocation rate is ~1.2GB/sec across all measurements


Sample responsiveness improvement



- SpecJBB + Slow churning 2GB LRU Cache
- Live set is ~2.5GB across all measurements
- Allocation rate is ~1.2GB/sec across all measurements

Instance capacity test: "Fat Portal"


CMS: Peaks at ~ 3GB / 45 concurrent users



* LifeRay portal on JBoss @ 99.9% SLA of 5 second response times

Instance capacity test: "Fat Portal"

C4: still smooth @ 800 concurrent users

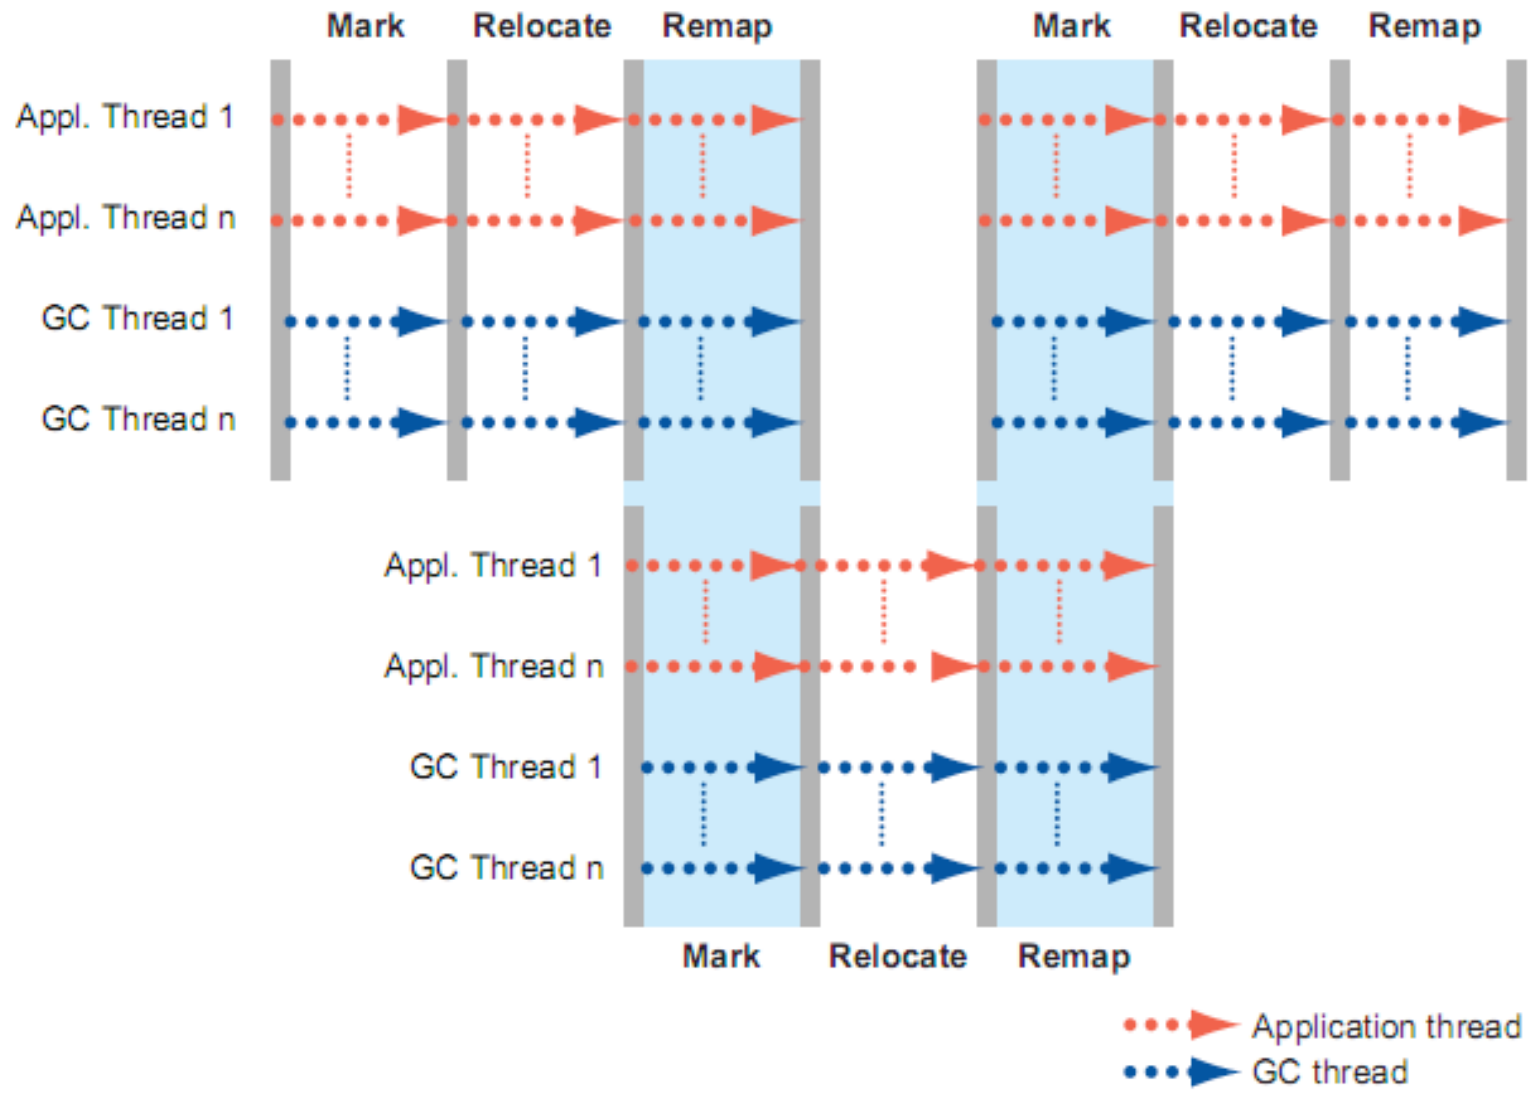


C4 Algorithm fundamentals

C4 algorithm highlights

- ☉ Same core mechanism used for both generations
 - ☉ Concurrent Mark-Compact
- ☉ A Loaded Value Barrier (LVB) is central to the algorithm
 - ☉ Every heap reference is verified as “sane” when loaded
 - ☉ “Non-sane” refs are caught and fixed in a self-healing barrier
- ☉ Refs that have not yet been “marked through” are caught
 - ☉ Guaranteed single pass concurrent marker
- ☉ Refs that point to relocated objects are caught
 - ☉ Lazily (and concurrently) remap refs, no hurry
 - ☉ Relocation and remapping are both concurrent
- ☉ Uses “quick release” to recycle memory
 - ☉ Forwarding information is kept outside of object pages
 - ☉ Physical memory released immediately upon relocation
 - ☉ “Hand-over-hand” compaction without requiring empty memory

The C4 GC Cycle

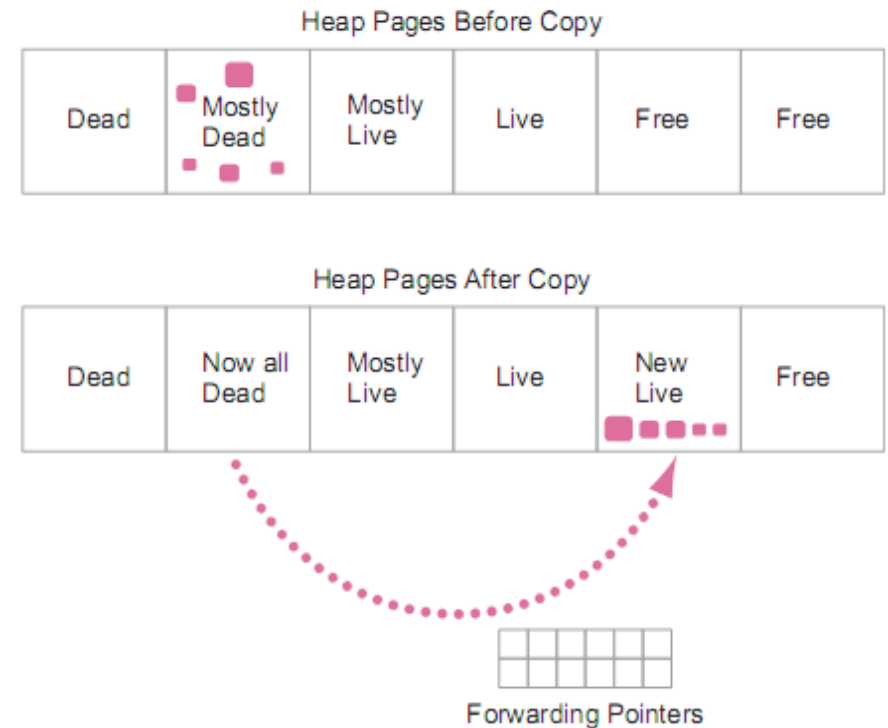


Mark Phase

- ④ Mark phase finds all live objects in the Java heap
- ④ Concurrent, predictable: always complete in a single pass
- ④ Uses LVB to defeat concurrent marking races
 - ④ Tracks object references that have been traversed by using an “NMT” (not marked through) metadata state in each object reference
 - ④ Any access to a not-yet-traversed reference will trigger the LVB
 - ④ Triggered references are queued on collector work lists, and reference NMT state is corrected
 - ④ “Self healing” corrects the memory location that the reference was loaded from
- ④ Marker tracks total live memory in each memory page
 - ④ Compaction uses this information to go after the sparse pages first
(But each cycle will tend to compact the entire heap...)

Relocate Phase

- Compacts to reclaim heap space occupied by dead objects in “from” pages without stopping mutator
- Protects “from” pages.
- Uses LVB to support concurrent relocation and lazy remapping by triggering on any access to references to “from” pages
- Relocates any live objects to newly allocated “to” pages
- Maintains forwarding pointers outside of “from” pages
- Virtual “from” space cannot be recycled until all references to relocated objects are remapped




“Quick Release”: Physical memory can be immediately reclaimed, and used to feed further compaction or allocation

Remap Phase

- ④ Scans all live objects in the heap
- ④ Looks for references to previously relocated objects, and updates (“remaps”) them to point to the new object locations
- ④ Uses LVB to support lazy remapping
 - ④ Any access to a not-yet-remapped reference will trigger the LVB
 - ④ Triggered references are corrected to point to the object’s new location by consulting forwarding pointers
 - ④ “Self healing” corrects the memory location the reference was loaded from
- ④ Overlaps with the next mark phase’s live object scan
 - ④ Mark & Remap are executed as a single pass

The C4 GC Cycle



GC Tuning

Java GC tuning is "hard"...

Examples of actual command line GC tuning parameters:

```
Java -Xmx12g -XX:MaxPermSize=64M -XX:PermSize=32M -XX:MaxNewSize=2g  
-XX:NewSize=1g -XX:SurvivorRatio=128 -XX:+UseParNewGC  
-XX:+UseConcMarkSweepGC -XX:MaxTenuringThreshold=0  
-XX:CMSInitiatingOccupancyFraction=60 -XX:+CMSScavengeBeforeRemark  
-XX:+UseCMSInitiatingOccupancyOnly -XX:ParallelGCThreads=12  
-XX:LargePageSizeInBytes=256m ...
```

```
Java -Xms8g -Xmx8g -Xmn2g -XX:PermSize=64M -XX:MaxPermSize=256M  
-XX:-OmitStackTraceInFastThrow -XX:SurvivorRatio=2 -XX:-UseAdaptiveSizePolicy  
-XX:+UseConcMarkSweepGC -XX:+CMSConcurrentMTEnabled  
-XX:+CMSScavengeBeforeRemark -XX:+CMSScavengeBeforeSurvivorRemark  
-XX:CMSMaxAbortablePreCleanTime=10000 -XX:+UseCMSInitiatingOccupancyOnly  
-XX:CMSInitiatingOccupancyFraction=63 -XX:+UseParNewGC -Xnoclassgc ...
```

The complete guide to Zing GC tuning

`java -Xmx40g`

Q & A

G. Tene, B. Iyengar and M. Wolf

C4: The Continuously Concurrent Compacting Collector

In Proceedings of the international symposium on Memory management, ISMM'11, ACM, pages 79-88

Jones, Richard; Hosking, Antony; Moss, Eliot (25 July 2011).

The Garbage Collection Handbook: The Art of Automatic Memory Management. CRC Press. ISBN 1420082795.